

# BFT Protocols

Introduction to Blockchain Science and Engineering

Aggelos Kiayias

Nikos Leonardos, Dionysis Zindros, Christos Nasikas

# Byzantine Consensus (Binary inputs)

$n$  parties  $(1, 2, \dots, n)$ ,  $t$  adversarial.

Let  $x_i \in \{0, 1\}$  be the input of party  $i$ .

Honest parties should *decide* on values  $y_i \in \{0, 1\}$  satisfying the following two properties.

- **Agreement:** if parties  $i$  and  $j$  are honest, then  $y_i = y_j$ .
- **Validity:** if there exists  $v \in \{0, 1\}$  such that  $x_i = v$  for each honest party  $i$ , then  $y_i = v$  for each honest party  $i$ .

Furthermore, they reach a decision in finite time.

# Exponential Information Gathering Algorithm (EIG)


## Algorithm Sketch.

- At round 1, send to everyone your input.
- At round  $r+1$ , send to everyone all messages you received at round  $r$  (avoiding redundant messages).

Each party arranges the messages in each own EIG tree.

- Let  $u_1, \dots, u_n$  be the messages received in the first round.
- Subsequently,  $u_{xj}$  is the value received from  $j$  as the value  $u_x$  in  $j$ 's tree.

Note: there need be no repetitions in the label of a node (e.g.,  $x$  in  $u_x$  should contain distinct identifiers).



The value party 2 told me that party 3 send him in the previous round.

# EIG Termination

The EIG algorithm terminates after  $t+1$  rounds. The output value of each party is defined as follows.

- For each leaf  $v$  in the EIG tree, set  $z_v = u_v$ .
- For an internal node  $v$ , set  $z_v$  equal to the majority of the  $z$ -values of its children. If the majority is not defined, set  $z_v = z_0$ , for some default value  $z_0$ .
- Define the output as  $z_{\text{root}}$ .

# Impossibility results I

**Theorem[LSP1982]** Impossible for  $n < 3t + 1$ .

**Theorem[FL1982]** Impossible in  $t$  rounds.

**Example** The EIG algorithm with  $t=1$  needs at least 2 rounds.

1. If a party received a single 1, its output should be 1. (Because the 1 could be coming from the adversary.)
2. If a party received two 1s, its output should be 1. (Because one of them could have been sent from the adversary, while another party received a single 1 and will decide on 1 according to the previous statement.)
3. And so on...

**Theorem[GM1998]** Doable for  $n > 3t$  in  $t+1$  rounds.

# Impossibility results II

**Theorem[BT1985]** Asynchronous Byzantine Consensus is impossible with  $n < 3t + 1$ , even if the parties have agreed on a PKI.

**Proof** Partition parties into sets A, B, C of size at most  $t$ . Consider 3 scenarios.

- A. A malicious, B and C honest with inputs 0. The adversary sends no messages. The honest parties should decide on 0 until some time  $T_A$ .
- B. B malicious, A and C honest with inputs 1. The adversary sends no messages. The honest parties should decide on 1 until some time  $T_B$ .
- C. C malicious, B and A honest with inputs 0 and 1 respectively. The adversary communicates with B as the honest C in scenario A and with A as the honest C in scenario B. At the same time every communication between A and B is delayed for time at least  $\max\{T_A, T_B\}$ .

The crux is that A has the same view in scenarios B and C. Similarly for B, in scenarios A and C. Agreement in scenario C is impossible, if validity is achieved in scenarios A and B.

# A blockchain related to proof-of-stake

Servers  $S_1, \dots, S_n$  with shared verification keys  $pk_1, \dots, pk_n$  and private signing keys  $sk_1, \dots, sk_n$ .

$B_0$ : Genesis block containing the public info.

$B_i = (k, d, sl, \sigma_{sl}, \sigma_{block})$ , where

$k$ : hash of previous block

$d$ : data

$sl$ : slot number


$\sigma_{sl}$ : signature on  $sl$  by  $S_{sl \bmod n}$

$\sigma_{block}$ : signature on whole block by  $S_{sl \bmod n}$

# Characteristic sequences and executions

The characteristic sequence of an execution with  $L$  slots  $0, 1, \dots, L-1$  is a binary string  $w$  in  $\{0, 1\}^L$  such that  $w_i = 1$  iff  $S_{i \bmod n}$  is adversarial.

**Example**  $w = 000111000$ . Squares denote honest slots (including the genesis) and circles the adversarial ones. The adversary keeps the chain ending with slot 5 on top hidden so that  $S_6$  extends the bottom chain. This results in a **fork**, two disjoint chains (except for the genesis block) of maximum length. The corresponding characteristic sequence  $w$  is called **forkable**.





# Forkable sequences

**Fact** If  $\text{weight}(w) < \text{length}(w)/3$ , then  $w$  is not forkable.

**Proof** Let  $n = \text{length}(w)$  and  $t = \text{weight}(w)$  and assume  $w$  is forkable. The two chains must have length at least  $n-t$  (prove formally by induction). Thus,  $2(n-t)$  is a lower bound on the sum of their lengths. On the other hand, the  $n-t$  honest parties have contributed at most  $n-t$ , while the adversarial parties have contributed at most  $2t$  (because a given chain contains a slot at most once). Thus,  $(n-t) + 2t$  is an upper bound on the sum of their lengths. We have

$$2(n-t) \leq (n-t) + 2t \implies n \leq 3t.$$

# Consensus inspired from proof-of-stake [KR2018]

- Fix an arbitrary ordering of the servers:  $S_1, \dots, S_n$ .
- Construct a blockchain for  $5t+2$  rounds, recording your own input bit as data in any block you create.
- Upon termination output the majority of the first  $2t+1$  blocks of your chain.

**Theorem** If  $n > 3t$ , the protocol satisfies agreement and validity.

**Proof [KR2018]** It can be shown that the first  $2t+1$  blocks are common to all honest parties; this implies agreement. Validity follows from the fact that among the first  $2t+1$  at most  $t$  are adversarial and so the majority of them belong to honest parties.

# Bitcoin Consensus

- Miners run the Bitcoin protocol recording their own input bit as data in any block they compute.
- When their chain has at least  $2k$  blocks (for some security parameter  $k$ ), the broadcast it and stop.
- Output is the majority of the bits recorded in the first  $k$  blocks.

**Theorem [GKL15]** If  $t < n/3$ , the above protocol satisfies Agreement and Validity with probability  $1 - e^{-\Omega(k)}$ .

**Remark** For Agreement,  $t < n/2$  suffices.

# Common-Prefix Property and Agreement

**Common-Prefix Property** For any pair of honest parties adopting the chains  $C_1$  and  $C_2$  at rounds  $r_1$  and  $r_2$  respectively. If  $r_1 \leq r_2$ , then  $C_1[-k]$  is a prefix of  $C_2$ , where  $C_1[-k]$  is  $C_1$  without its last  $k$  blocks.

Common-Prefix Property implies Agreement. This is because the parties are pruning at least  $k$  blocks from their chains when keeping only their initial  $k$  blocks. Thus, the initial  $k$  blocks are common to all honest parties.

In [GKL2015] it is shown that in Bitcoin the Common-Prefix fails with probability exponentially small in  $k$ , if the adversary's hashing power is sufficiently bounded below  $1/2$ . It follows that Agreement is satisfied with probability  $1 - e^{-\Omega(k)}$ , when  $t$  is sufficiently less than  $n/2$ .

# Chain-Quality Property and Validity

**Chain-Quality Property (informal)** Among any sufficiently large number of consecutive blocks in an honest party's chain, a fraction of at least  $(n-2t)/(n-t)$  have been computed by honest parties.

Chain-Quality implies Validity, when  $t$  is sufficiently less than  $n/3$ . This is because, in that case, the majority of the first  $k$  blocks have been computed by honest parties. Thus, if all honest parties have input  $v$ , the majority of values recorded in the first  $k$  blocks will be  $v$ .

In [GKL2015] it is shown that in Bitcoin Chain-Quality fails with probability exponentially small in  $k$ , if the adversary's hashing power is sufficiently bounded below  $1/2$ . It follows that Validity is satisfied with probability  $1 - e^{-\Omega(k)}$ , when  $t$  is sufficiently less than  $n/3$ .

# References

**[BT1985]** Bracha, Toueg. Asynchronous consensus and broadcast protocols.

**[FL1982]** Fischer, Lynch. A lower bound on the time to assure interactive consistency.

**[GKL2015]** Garay, Kiayias, Leonardos. The Bitcoin Backbone Protocol: Analysis and Applications.

**[GM1998]** Garay, Moses. Fully polynomial Byzantine agreement for  $n > 3t$  processors in  $t+1$  rounds.

**[KR2018]** Kiayias, Russell. Ouroboros-BFT. A simple Byzantine Fault Tolerant Consensus Protocol.

**[LSP1982]** Lamport, Shostak, Pease. The Byzantine generals problem.